

# Machine Learning

Permutation importance with RandomForestClassifier and DecisionTreeClassifier

Kunal Khurana

2024-01-13

# Table of contents

eli5 library for permutation importance . . . . .	4
Practice . . . . .	4
Hypothesis testing . . . . .	7

- **Process-**

get a trained model

calculate the loss function

return the data into original orde

```
#import eli5
#from eli5.sklearn import PermutationImportance

from sklearn.datasets import load_iris
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
```

```
data = pd.read_csv('FIFA 2018 Statistics.csv')
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 128 entries, 0 to 127
```

```
Data columns (total 27 columns):
```

#	Column	Non-Null Count	Dtype
0	Date	128 non-null	object
1	Team	128 non-null	object
2	Opponent	128 non-null	object
3	Goal Scored	128 non-null	int64
4	Ball Possession %	128 non-null	int64

```

5   Attempts                128 non-null    int64
6   On-Target                128 non-null    int64
7   Off-Target               128 non-null    int64
8   Blocked                  128 non-null    int64
9   Corners                  128 non-null    int64
10  Offsides                  128 non-null    int64
11  Free Kicks                128 non-null    int64
12  Saves                     128 non-null    int64
13  Pass Accuracy %           128 non-null    int64
14  Passes                    128 non-null    int64
15  Distance Covered (Kms)    128 non-null    int64
16  Fouls Committed           128 non-null    int64
17  Yellow Card               128 non-null    int64
18  Yellow & Red              128 non-null    int64
19  Red                       128 non-null    int64
20  Man of the Match          128 non-null    object
21  1st Goal                  94 non-null     float64
22  Round                     128 non-null    object
23  PSO                       128 non-null    object
24  Goals in PSO              128 non-null    int64
25  Own goals                  12 non-null     float64
26  Own goal Time             12 non-null     float64

```

dtypes: float64(3), int64(18), object(6)

memory usage: 27.1+ KB

```
data.head()
```

	Date	Team	Opponent	Goal Scored	Ball Possession %	Attempts	On-Target	Off
0	14-06-2018	Russia	Saudi Arabia	5	40	13	7	3
1	14-06-2018	Saudi Arabia	Russia	0	60	6	0	3
2	15-06-2018	Egypt	Uruguay	0	43	8	3	3
3	15-06-2018	Uruguay	Egypt	1	57	14	4	6
4	15-06-2018	Morocco	Iran	0	64	13	3	6

```

#
y = (data['Man of the Match'] == 'Yes')
feature_names = [i for i in data.columns if data[i].dtype in [np.int64]] #choosing integers
X = data[feature_names]
train_X, val_X, train_y, val_y = train_test_split(X, y, random_state = 1)

```

```
my_model = RandomForestClassifier(n_estimators = 100,
                                random_state = 0).fit(train_X, train_y)
```

## eli5 library for permutation importance

### Practice

```
import pandas as pd
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
```

```
import seaborn as sns
import matplotlib.pyplot as plt
```

```
data2= pd.read_csv('train.csv')
```

```
data2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   key                   50000 non-null  object
1   fare_amount           50000 non-null  float64
2   pickup_datetime      50000 non-null  object
3   pickup_longitude     50000 non-null  float64
4   pickup_latitude      50000 non-null  float64
5   dropoff_longitude    50000 non-null  float64
6   dropoff_latitude     50000 non-null  float64
7   passenger_count      50000 non-null  int64
dtypes: float64(5), int64(1), object(2)
memory usage: 3.1+ MB
```

```
data2.columns
```

```
Index(['key', 'fare_amount', 'pickup_datetime', 'pickup_longitude',
      'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude',
      'passenger_count'],
      dtype='object')
```

```
data2.head()
```

	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	2009-06-15 17:26:21.00000001	4.5	2009-06-15 17:26:21 UTC	-73.844311	40.721319
1	2010-01-05 16:52:16.00000002	16.9	2010-01-05 16:52:16 UTC	-74.016048	40.711303
2	2011-08-18 00:35:00.000000049	5.7	2011-08-18 00:35:00 UTC	-73.982738	40.761270
3	2012-04-21 04:30:42.00000001	7.7	2012-04-21 04:30:42 UTC	-73.987130	40.733143
4	2010-03-09 07:51:00.000000135	5.3	2010-03-09 07:51:00 UTC	-73.968095	40.768008

```
# Remove data with extreme outlier coordinates or negative fares
data2 = data2.query('pickup_latitude > 40.7 and pickup_latitude < 40.8 and ' +
                    'dropoff_latitude > 40.7 and dropoff_latitude < 40.8 and ' +
                    'pickup_longitude > -74 and pickup_longitude < -73.9 and ' +
                    'dropoff_longitude > -74 and dropoff_longitude < -73.9 and ' +
                    'fare_amount > 0'
                    )
```

```
# assigning prediction target
```

```
y = data2.fare_amount
```

```
base_features = ['pickup_longitude',
                 'pickup_latitude',
                 'dropoff_longitude',
                 'dropoff_latitude',
                 'passenger_count']
```

```
X = data2[base_features]
```

```
# train the model
```

```
train_X, val_X, train_y, val_y = train_test_split(X, y, random_state = 1)
first_model = RandomForestRegressor(n_estimators = 50, random_state = 1).fit(train_X, train_y)
```

```
train_X.describe()
```

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
count	37500.000000	37500.000000	37500.000000	37500.000000	37500.000000
mean	-72.520688	39.931410	-72.514256	39.930631	1.670827
std	10.360818	6.010290	10.378988	6.014321	1.293083
min	-75.414728	-74.006893	-84.654241	-74.006377	0.000000
25%	-73.992092	40.735008	-73.991196	40.734333	1.000000
50%	-73.981811	40.752825	-73.980012	40.753397	1.000000
75%	-73.967051	40.767428	-73.963292	40.768243	2.000000
max	40.783472	42.160275	40.802437	42.168717	6.000000

```
train_y.describe()
```

```
count    37500.000000
mean         11.402322
std          9.715174
min         -5.000000
25%          6.000000
50%          8.500000
75%         12.500000
max         180.000000
Name: fare_amount, dtype: float64
```

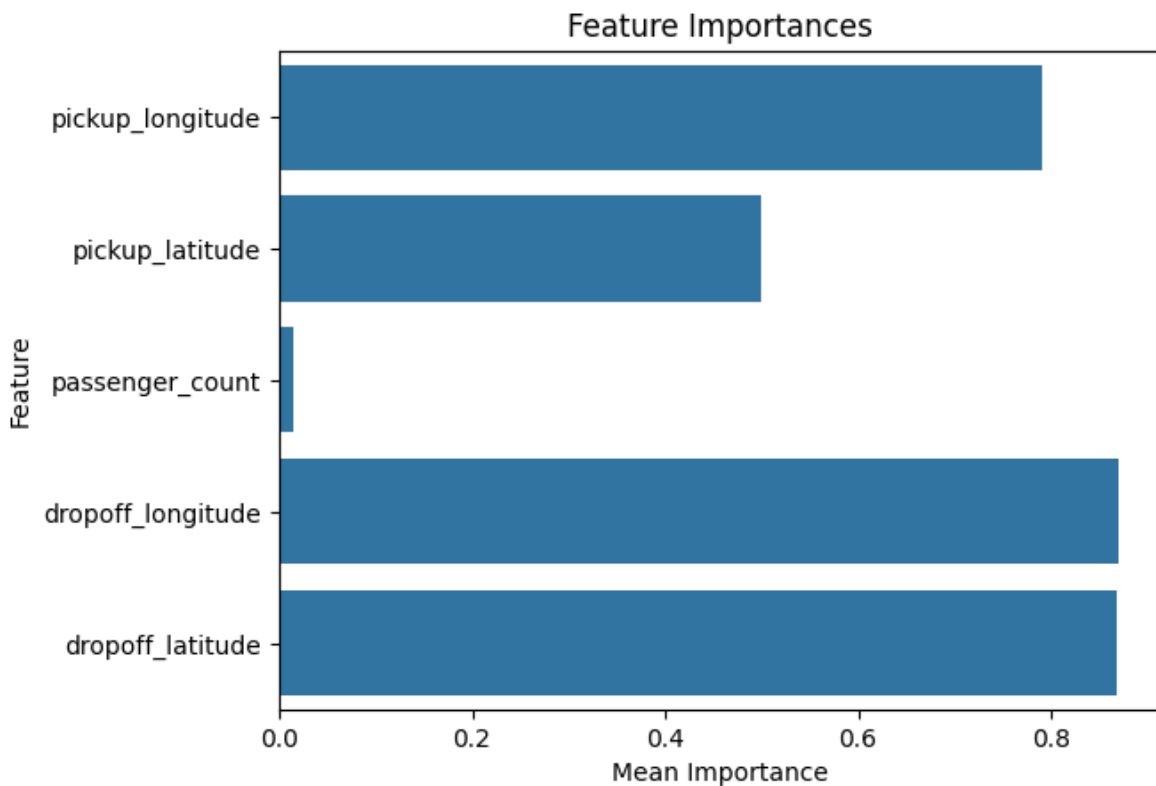
```
train_X, val_X, train_y, val_y = train_test_split(X, y, random_state = 1)
first_model = RandomForestRegressor(n_estimators = 50, random_state = 1).fit(train_X, train_y)
```

```
from sklearn.inspection import permutation_importance
```

```
result = permutation_importance(first_model, train_X, train_y, n_repeats=30, random_state=1)
```

```
# resulting importances
importance_df = pd.DataFrame({'feature': base_features,
                             'importance_mean': result.importances_mean,
                             'importance_std' : result.importances_std})
importance_df = importance_df.sort_values('feature',
                                         ascending = False)
```

```
# visualize
sns.barplot(x = 'importance_mean', y = 'feature', data= importance_df)
plt.title('Feature Importances')
plt.xlabel('Mean Importance')
plt.ylabel('Feature')
plt.show()
```



## Hypothesis testing

Why are latitudes have larger effect?

1. Travel might tend to have greater latitude distances than longitude distances.
2. Different parts of the city might have different pricing rules (e.g. price per mile), and pricing rules could vary more by latitude than longitude.
3. Tolls might be greater on roads going North<->South (changing latitude) than on roads going East <-> West (changing longitude). Thus latitude would have a larger effect on

the prediction because it captures the amount of the tolls.

```
# reducing length of X
X_trimmed = X.iloc[:len(y)]

# create new features

data2['abs_lon_change'] = abs(data2.dropoff_longitude - data2.pickup_longitude)
data2['abs_lat_change'] = abs(data2.dropoff_latitude - data2.pickup_latitude)

features_2 = ['pickup_longitude',
              'pickup_latitude',
              'dropoff_longitude',
              'dropoff_latitude',
              'passenger_count',
              'abs_lon_change',
              'abs_lat_change']
X= data2[features_2]


new_train_X, new_val_X, new_train_y, new_val_y = train_test_split(X, y, random_state = 1)
second_model = RandomForestRegressor(n_estimators = 50, random_state = 1).fit(new_train_X,

result2 = permutation_importance(second_model, new_train_X, new_train_y, n_repeats=30, ran

# resulting importances
importance_df2 = pd.DataFrame({'feature': features_2,
                              'importance_mean': result2.importances_mean,
                              'importance_std' : result2.importances_std})
importance_df2 = importance_df2.sort_values('feature',
                                           ascending = False)

# visualize
sns.barplot(x = 'importance_mean', y = 'feature', data= importance_df2)
plt.title('Feature Importances')
plt.xlabel('Mean Importance')
plt.ylabel('Feature')
plt.show()
```





```
X.describe()
```

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	abs_
count	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000
mean	-72.509756	39.933759	-72.504616	39.926251	1.667840	0.151
std	10.393860	6.224857	10.407570	6.014737	1.289195	3.077
min	-75.423848	-74.006893	-84.654241	-74.006377	0.000000	0.000
25%	-73.992062	40.734880	-73.991152	40.734372	1.000000	0.005
50%	-73.981840	40.752678	-73.980082	40.753372	1.000000	0.012
75%	-73.967148	40.767360	-73.963584	40.768167	2.000000	0.023
max	40.783472	401.083332	40.851027	43.415190	6.000000	74.01

```
print(X.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
```

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	pickup_longitude	50000 non-null	float64
1	pickup_latitude	50000 non-null	float64
2	dropoff_longitude	50000 non-null	float64
3	dropoff_latitude	50000 non-null	float64
4	passenger_count	50000 non-null	int64
5	abs_lon_change	50000 non-null	float64
6	abs_lat_change	50000 non-null	float64

dtypes: float64(6), int64(1)

memory usage: 2.7 MB

None

```
print(y.info())
```

<class 'pandas.core.series.Series'>

Index: 31289 entries, 2 to 49999

Series name: fare\_amount

Non-Null Count	Dtype
----------------	-------

31289 non-null	float64
----------------	---------

dtypes: float64(1)

memory usage: 488.9 KB

None